

Software vulnerabilities: memory corruption

Classifying vulnerabilities

- Memory corruption (safety)
 - Buffer Overflow
 - Integer Overflow
 - Stack Overflow
 - Use after free/ Double free
 - Null pointer dereference
- Exploiting memory corruption to control program
 - Format string attacks
 - Stack buffer overflow (or stack-based buffer overflow)
 - Heap buffer overflow (or heap-based buffer overflow)
 - Overflowing on non “control-data” memory
- Resulting attacks
 - Code injection on the stack (e.g., shellcodes)
 - Code reuse techniques (Return-to-libc, Return-Oriented Programming, Jump-Oriented Programming)


Buffer Overflows

- A buffer overflow occurs any time a program attempts to store data beyond the boundaries of a buffer, overwriting the adjacent memory locations
- Originates from mistakes done while writing code
 - unfamiliarity with language
 - Boundary or arithmetic errors
- Mostly C / C++ programs
- Addressed by languages with automatic memory management: dynamic bounds checks (e.g., Java) or automatic resizing of buffers (e.g., Perl)
 - Still native libraries that are written in C (e.g., JNI)

Buffer Overflows: Basics

```
char src[]="ABCDEFGHI";
char tmp_buff[5];
int password_checked;


strcpy(tmp_buff, src);
```



Buffer Overflows: After strcpy()

```
char src[]="ABCDEFGHI";
char tmp_buff[5];
int password_checked;

strcpy(tmp_buff, src);
```



Buffer Overflows: usual suspects

- String manipulation functions that don't properly check string length
 - gets(), strcpy(), ...
- Copy with incorrect parameters
 - memcpy() ...
- Incorrect computation of required memory length
 - zero-sized malloc()s ...

Preventing Buffer Overflows


- Safe usage of copy primitives using libc methods
 - For instance, limit amount of data to copy
 - `strncpy(tmp_buff, src, sizeof(tmp_buf))`
- Compiler extension if array length can be computed
- Annotations (e.g. Deputy, CCured...)

Integer overflows

- Unsigned ints
 - 32 bits: range from 0 to $2^{32}-1$ (4,294,967,295)
 - Overflow: $4,294,967,295 + 1 = 0$
- Signed ints
 - 32 bits: range from $-(2^{31})$ to $2^{31} - 1$ (2,147,483,647)
 - Overflow: $2,147,483,647+2 = -2,147,483,648$
- Some languages (Java, Ada) throw exceptions, many don't
- Attacker can supply large values used in :
 - Computation of the size of a buffer allocation (malloc)
 - Array access (in particular the bound checks)
- Solution: always check overflow before critical operation!

More about Memory Layout

- Text
 - Also called code segment
- Data
 - Global initialized data
- BSS
 - Global uninitialized data
- Stack
 - Local variables
 - Also used to store function environments and parameters during calls (stack frames)
 - LIFO
 - Multi-threading : multiple stacks
- Heap
 - Dynamically allocated variables
 - Reserved through `calloc()` and `malloc()`



Program Memory Stack


```
int x = 100;
int main()
{
    // data stored on stack
    int a=2;
    float b=2.5;
    static int y;

    // allocate memory on heap
    int *ptr = (int *) malloc(2*sizeof(int));

    // values 5 and 6 stored on heap
    ptr[0]=5;
    ptr[1]=6;

    // deallocate memory on heap
    free(ptr);

    return 1;
}
```



Order of the function arguments in stack


```
void func(int a, int b)
{
    int x, y;

    x = a + b;
    y = a - b;
}
```



```
movl 12(%ebp), %eax      ; b is stored in %ebp + 12
movl 8(%ebp), %edx       ; a is stored in %ebp + 8
addl %edx, %eax
movl %eax, -8(%ebp)      ; x is stored in %ebp - 8
```

Function Call Stack

```
void f(int a, int b)
{
    int x;
}
void main()
{
    f(1,2);
    printf("hello world");
}
```



Stack Layout for Function Call Chain



Stack Overflow

- Stack Overflows can be caused by:
 - Recursive calls
 - Reentrant interrupts
 - Allocations on the stack
 - Large
 - Controlled by the attacker
 - `Alloca()`, `char array[function_parameter]`

Stack Overflow

- This should never happen ?
- Detection : guard page (Unix)
 - Problem when allocation is too large
 - the stack pointer can “jump over” the guard page
- Prevention: Not so easy
 - Abstract interpretation [*Regehr05*]
 - Works as long as control flow can be determined statically
 - Forbid / bound recursion
 - Forbid / limit allocation on the stack
- Problem with micro-controllers
 - No MMU (Memory Management Unit)

Stack Buffer Overflow / Overrun

- A special case of buffer overflow
 - Exploitation of existing vulnerability in order to modify control flow
- The overflowed buffer is allocated on the stack, typically corrupts return address
- First wide-scale exploitation: Morris' worm (1989) in Unix's fingerd

Vulnerable Program

```
int main(int argc, char **argv)
{
    char str[400];
    FILE *badfile;

    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 300, badfile);
    foo(str); ←
    printf("Returned Properly\n");
    return 1;
}
```

- Reading 300 bytes of data from badfile.
- Storing the file contents into a str variable of size 400 bytes.
- Calling foo function with str as an argument.

Note : Badfile is created by the user and hence the contents are in control of the user.


Vulnerable Program

```
/* stack.c */
/* This program has a buffer overflow vulnerability. */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str)
{
    char buffer[100];

    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);

    return 1;
}
```




Consequences of Buffer Overflow

Overwriting return address with some random address can point to :



- Invalid instruction
- Non-existing address
- Access violation
- Attacker's code → Malicious code to gain access

How to Run Malicious Code



More on Modifying return address

- Not always easy to know the exact address
- Trampolines :
 - `jmp %esp` at a fixed address
- NOP sled
 - A long sequence of NOP
 - Followed by the shellcode
 - Jumping anywhere in there leads to the shellcode




Exploiting Stack Buffer Overflows

- Corrupting the Control Flow
 - Return address in a stack frame
 - Function pointer
 - Exception handler
 - Global Offset Table (GOT)
 - C++ objects
- Corrupting the Program State (other stack frames)
 - Variable that stores authentication status
 - Modify stored password
 - Pointer value

Protections: Canaries

- Objective : Detect unexpected modifications of values on the stack (e.g., return address)
- Inserting a known value on the stack
 - the “canary”
- Compiler extension - inserts code that:
 - Adds a random value after the return address
 - Checks canary value before using return address
- Limitations
 - Canaries can be guessed or obtained through memory leaks
 - Canary copy needs to be stored in secure place that might also be corrupted



Protections: Non Executable Memories (Writable XOR Executable / NX)

- Executable and writable memories regions
 - Allows to directly write instructions (in a buffer) and execute them, e.g., from the stack.
 - Makes stack-based buffer overflows easy to exploit
- W XOR X (A.k.a NX / XN / DEP / PAX)
 - OS splits memory regions – no region is both writable and executable
- Most systems now support a form of W XOR X:
 - Was complex on x86 (segmentation)
 - HW support for it since ~2004
- Many names for the same feature :
 - AMD : Enhanced Virus Protection,
 - ARM : XN eXecute Never
- Defeated by Return-to-libc attacks / Return Oriented programming

Protections: ASLR (Address Space Layout Randomization)

- Randomizes base addresses of memory segments
 - Addresses change at every execution
- Randomize start or base address of:
 - Program code
 - Libraries code
 - Heap/Stack/Data regions
- Objectives:
 - Difficult to guess the stack address in the memory.
 - Difficult to guess %ebp address and address of the malicious code
- Many programs have problems with that
 - Sometimes rewriting part of them is necessary, notably when they contain specially crafted assembly code
 - Code needs to be position independent for program randomization
- Limitations
 - Memory leaks used to learn memory layout
 - Address space / system limitations (e.g., page boundaries) may allow brute-force probes

Address Space Layout Randomization

```
#include <stdio.h>
#include <stdlib.h>

void main()
{
    char x[12];
    char *y = malloc(sizeof(char)*12);

    printf("Address of buffer x (on stack) : 0x%x\n", x);
    printf("Address of buffer y (on heap) : 0x%x\n", y);
}
```

Address Space Layout Randomization : Working

```
$ sudo sysctl -w kernel.randomize_va_space=0  
kernel.randomize_va_space = 0  
$ a.out  
Address of buffer x (on stack): 0xbffff370  
Address of buffer y (on heap) : 0x804b008  
$ a.out  
Address of buffer x (on stack): 0xbffff370  
Address of buffer y (on heap) : 0x804b008
```

1

```
$ sudo sysctl -w kernel.randomize_va_space=1  
kernel.randomize_va_space = 1  
$ a.out  
Address of buffer x (on stack): 0xbf9deb10  
Address of buffer y (on heap) : 0x804b008  
$ a.out  
Address of buffer x (on stack): 0xbf8c49d0  
Address of buffer y (on heap) : 0x804b008
```

2

```
$ sudo sysctl -w kernel.randomize_va_space=2  
kernel.randomize_va_space = 2  
$ a.out  
Address of buffer x (on stack): 0xbf9c76f0  
Address of buffer y (on heap) : 0x87e6008  
$ a.out  
Address of buffer x (on stack): 0xbfe69700  
Address of buffer y (on heap) : 0xa020008
```

3